

Low-Latency Cross-Venue Arbitrage: Binance → Polymarket

Strategy, Infrastructure, Deployment, and a Production Post-Mortem

Cullen Trace Campana

Contents

1	Executive Summary	3
1.1	The Trade Idea	3
1.2	The System	3
1.3	What Was Validated and What Was Not	4
1.4	Lessons	4
2	The Trade Idea	5
2.1	Polymarket Short-Dated Crypto Markets	5
2.2	Where the Edge Comes From	5
2.3	The Edge, Measured	5
2.4	Why the Edge Is Structurally Thin	6
3	Signal Detection (Tokyo)	7
3.1	The Three-Gate Detector	7
3.2	Confidence Scoring	7
3.3	Per-Asset Calibration	7
3.4	A Meta-Labeling Layer	8
4	Execution Logic (Dublin)	9
4.1	Entry Gates	9
4.2	Order Execution	9
4.3	Exit Paths	10
4.4	Sizing and Risk Limits	10
5	Infrastructure	11
5.1	Two-Node Topology	11
5.2	The Wire Protocol	11
5.3	Node-Internal Design	12
5.4	Mock-to-Live Parity	12
5.5	Latency Budget	13
6	Deployment	14
6.1	Instances and Placement	14
6.2	Measured Latencies	14
6.3	Operations	14
7	Backtest Performance	15
7.1	Headline Results	15
7.2	Fees Dominate	15
7.3	The Edge Has Structure	16
7.4	Why the Tick-Rate Effect: Reversal, Not Throughput	17
8	Production Results and Post-Mortem	19
8.1	What Happened	19
8.2	Root Cause: Spoofed Liquidity	19
8.3	Why Latency Was Not the Problem	19
8.4	Adverse Selection, Restated	19

9	Lessons and What I Would Do Differently	20
9.1	Validate Liquidity, Not Just Signal	20
9.2	Model Cancellation-on-Arrival	20
9.3	Respect the Geography	20
9.4	Where This Approach Does Work	20
10	Conclusion	21
A	System Reference	22
A.1	Component Map	22
A.2	Backtest Summary	23
A.3	Key Parameters	23

1 Executive Summary

This report documents a low-latency trading system built to exploit a specific, well-defined inefficiency: Polymarket’s short-dated crypto “up-or-down” prediction markets reprice with a measurable lag relative to the underlying spot price on Binance. The system was engineered to production depth: two co-located compute nodes in different AWS regions, a hand-rolled 20-byte UDP signalling protocol, lock-free per-asset processing pipelines, a full mock-to-live execution harness, and per-asset parameter calibration against replayed market data. It was deployed to AWS and traded with real capital.

In replay, the strategy performed well. Across 8,576 trades drawn from roughly 52 hours of recorded order flow, it won 61.4% of the time with a positively-skewed per-trade P&L distribution and a stable upward equity curve. The signal was predictive, the microstructure edge was measurable and decayed as the thesis predicted, and the performance had sensible internal structure: it was strongest in quiet books, at prices away from 0.50, and at short holding periods, and it weakened in each of the opposite conditions.

It did not survive contact with the live market, and the reason it failed is the main finding of this report. The latency edge was real and was captured as designed. What broke the strategy was not speed, model quality, or engineering: it was that the liquidity it was pricing against did not exist. The displayed depth on Polymarket’s order book was, in large part, spoofed. In paper trading, where fills were simulated against the *displayed* book, the system filled essentially every marketable order. In live trading, where fills were matched against *real* resting size, roughly one order in ten filled. The edge was priced into a book that disappeared as soon as a real order tried to cross it.

1.1 The Trade Idea

Polymarket runs recurring 5-minute and 15-minute markets on whether BTC, ETH, SOL, or XRP will be higher or lower at the window’s close. The fair value of the “Yes” (up) token is a monotone function of the current underlying price relative to the window’s reference level and the time remaining. When Binance spot moves, that fair value changes instantly, but the Polymarket order book does not: quotes are updated by market makers and arbitrageurs reacting to the same Binance feed, and that reaction takes an observed 150–500 ms. The trade is to detect a Binance move, decide which outcome token it favours, and lift the correct side before the book reprices.

1.2 The System

Detection and execution are geographically separated because the two venues are. Binance’s matching engine is best reached from AWS **ap-northeast-1** (Tokyo); Polymarket’s CLOB is best reached from AWS **eu-west-1** (Dublin). A **Tokyo** node watches the Binance book and trade streams, runs a three-gate move detector, and fires a compact UDP signal to Dublin. A **Dublin** node receives the signal, checks the live Polymarket book through a set of entry gates, and executes a fill-and-kill order through an ECDSA-signed CLOB client. The path from a Binance price move to an order arriving at Polymarket was measured at roughly 140–150 ms, dominated by the ~ 120 ms transatlantic hop between the two nodes.

1.3 What Was Validated and What Was Not

The *signal* was validated: recorded data shows Polymarket repricing toward the signal within the first second or two after a Binance move (Section 2.3), and the backtested strategy was profitable with coherent conditional structure (Section 7). The *latency budget* was validated: the system consistently delivered orders inside the repricing window. The *execution surface* was not, and could not have been, validated in paper, because the paper harness inherited the same blind spot that the market exploited: it trusted the displayed book. This is the central lesson (Section 8).

1.4 Lessons

Three conclusions carry beyond this specific strategy. First, in a latency arbitrage the binding constraint is often not your own speed but the *realisability* of the liquidity you are racing toward; displayed size is not the same as available size. Second, a paper-trading harness that simulates fills against the displayed book will systematically overstate the performance of any strategy that takes liquidity, because it cannot see cancellation-on-arrival. Third, geography is a hard constraint: when the price source and the execution venue are on opposite sides of the world, the transit cost dominates every micro-optimisation, and local engineering cannot recover it.

A note on the data in this report. Much of the quantitative material here was reconstructed from recorded feed sessions and a replay backtest (the `regime_trade_log` and `edge_signals` datasets, 8,576 trades and 4,892 signals respectively). All backtest results should be read as *paper* results under a displayed-liquidity fill assumption; the divergence between them and the live outcome is the subject of Section 8.

2 The Trade Idea

2.1 Polymarket Short-Dated Crypto Markets

Polymarket lists a continuous series of short-horizon binary markets on major cryptocurrencies. For each asset in {BTC, ETH, SOL, XRP} and each timeframe in {5 min, 15 min}, a market opens asking whether the asset will close the window above or below its opening reference. Each market resolves to a pair of outcome tokens (“Yes” for up, “No” for down) that trade on a central limit order book priced in the interval (0, 1) and settle at \$1 for the correct outcome and \$0 for the incorrect one. A new window rotates in every 5 or 15 minutes; the system discovers active windows by constructing the exact market slugs (`btc-updown-5m-{ts}`) and querying Polymarket’s Gamma API.

2.2 Where the Edge Comes From

Let P_t be the Binance mid price of the underlying, R the window’s reference level, and τ the time remaining. The fair value of the up token is approximately

$$V_t^{\text{up}} \approx \Pr[P_{t+\tau} > R | P_t],$$

which, for the short horizons in question, is a smooth increasing function of the current distance $P_t - R$. The key point is that V_t^{up} responds to P_t *instantaneously*, whereas the Polymarket order book responds only through the actions of participants who are themselves watching Binance and re-quoting. That reaction is not instantaneous. In this market it was observed to take on the order of 150–500 ms.

2.3 The Edge, Measured

The lag is visible in the recorded data. For every detected Binance move, the `edge_signals` dataset records the Polymarket mid at fixed horizons afterward and computes the repricing: how far the Polymarket price moved in the direction the signal predicted. Figure 1 shows the average. The book reprices toward the signal over the first ~500 ms to 2 s (median repricing half-life ≈ 786 ms), then the edge decays and, by 10 s, mean-reverts past its starting point. The window is real but short, and it is where a fast system must act.




Figure 1: Empirical Polymarket repricing (in cents, in the signalled direction) at fixed horizons after a Binance move, per asset and pooled. The edge is concentrated in the first ~ 2 seconds and reverts by 10 s, which motivates both the low-latency design and the short holding periods.

2.4 Why the Edge Is Structurally Thin

The edge is bounded on both sides. On one side, the repricing lag sets the maximum time the opportunity exists, and Figure 1 shows the useful part is barely two seconds. On the other, the physical distance between the price source (Binance, in Asia) and the execution venue (Polymarket, reachable from Europe or the US) imposes an ~ 120 ms one-way transit paid on every signal. The system therefore operates with little margin: it must detect, decide, transmit, re-decide, sign, and submit inside a window that the transit alone nearly consumes. This tension, a real but geographically expensive edge, motivated the two-node, minimal-protocol architecture of Section 5. Whether the remaining margin is enough is what a live deployment exists to answer.

3 Signal Detection (Tokyo)

3.1 The Three-Gate Detector

Detection runs in Tokyo, as physically close to Binance as AWS placement allows, so that a move is observed as early as possible. The detector consumes Binance `depth10@100ms` book snapshots and the `@trade` stream for all four assets, maintains 60 seconds of rolling per-asset metrics (mid, order-flow imbalance, VWAP, separated buy/sell volume, book imbalance, and a millisecond-resolution price ring), and evaluates a conjunction of three gates on every book update. A signal fires only when *all three* pass simultaneously:

- **Gate A (price move).** The Binance mid has moved at least a threshold number of basis points over a short lookback (25–150 ms). This is the cheapest test and is evaluated first, so that the $\sim 99\%$ of ticks with no meaningful move exit in a few hundred nanoseconds.
- **Gate B (tick density).** At least a threshold number of book events per 100 ms occurred in the recent window. Real moves arrive as bursts of activity; a slow drift on a thin book does not.
- **Gate C (dollar flow).** The signed traded dollar flow over a short window confirms the move’s direction and exceeds a threshold. Spoofed orders can push the mid and generate book events without anyone actually trading; requiring confirming *executed* flow filters book-manipulation moves.

A per-asset cooldown (roughly one Polymarket repricing cycle) debounces the detector so that a single sustained move produces one signal rather than a burst of them. On a no-signal tick the detector returns in ~ 400 ns; the full signal-to-wire path, including packing and transmission, was budgeted at 50–200 μ s.

3.2 Confidence Scoring

When all gates pass, the detector emits a confidence score in $[0, 1]$ used downstream for position sizing and as an entry gate on the Dublin side. The score is a weighted blend of how far each gate cleared its threshold, normalised so that 1.0 corresponds to roughly three times threshold on each axis:

$$c = \underbrace{0.5 \min\left(1, \frac{|\Delta_{\text{bps}}|}{3\theta_{\text{move}}}\right)}_{\text{move}} + \underbrace{0.3 \min\left(1, \frac{\text{flow}}{3\theta_{\text{flow}}}\right)}_{\text{flow}} + \underbrace{0.2 \min\left(1, \frac{\text{density}}{3\theta_{\text{td}}}\right)}_{\text{density}}.$$

The move term carries the most weight because the magnitude of the Binance move is the most direct predictor of how far the Polymarket fair value has decoupled from its displayed price.

3.3 Per-Asset Calibration

The four assets differ enough in tick size, volatility, and flow that a single threshold set is inappropriate. Every gate parameter is therefore calibrated per asset. Calibration is performed offline by replaying recorded Binance tick and trade data and using Optuna’s TPE sampler (150 trials per asset) to search the joint detector-and-exit parameter space. To keep each trial cheap, all rolling metrics are precomputed once at every candidate lookback window

(25–250 ms) so that a trial reduces to vectorised NumPy indexing plus a short cooldown pass. The objective rewards signals whose subsequent move persists in the signalled direction and penalises reversals (“fakeouts”). Table 1 shows representative calibrated values.

Table 1: Representative per-asset calibrated parameters (excerpt).

Asset	Move thr.	Move win.	Flow min.	Cooldown	Exit timeout
	(bps)	(ms)	(USD)	(s)	(s)
BTC	4.75	25	15,000	5.0	14
ETH	3.75	25	800	4.5	27
SOL	4.75	25	11,800	10.0	3
XRP	5.00	50	13,800	1.0	27

3.4 A Meta-Labeling Layer

Beyond the hand-tuned gates, a second research layer was built: per-asset, per-timeframe gradient-boosted (LightGBM) and random-forest meta-models trained to predict whether a raw detector signal would convert into a profitable trade, using the full microstructure feature vector at signal time (volatility ratios, order-flow imbalance, book depth, tick and trade rates, and the Polymarket spread and distance from 0.50). Filtering a fast, permissive detector with a slower, more selective model raises precision without slowing the hot path, since the model runs only on the ticks that already cleared the gates. Its findings are consistent with the conditional structure documented in Section 7.3.

4 Execution Logic (Dublin)

The Dublin node turns a directional signal into a filled position. Its design priority is the inverse of Tokyo’s: where Tokyo optimises for the earliest possible detection, Dublin optimises for *not trading* unless conditions are favourable, because on a thin edge a bad fill erases many good ones.

4.1 Entry Gates

Every incoming signal passes through ten entry gates, checked in order and short-circuiting on the first failure:

1. **Signal age.** Reject signals older than 500 ms (the edge is already gone).
2. **Clock skew.** Reject signals timestamped more than 50 ms in the future (clock desync or corruption).
3. **Confidence.** Reject signals below a minimum confidence (0.10).
4. **No-trade zone.** Refuse to open within 30 s of window expiry, where resolution risk dominates the repricing edge.
5. **Book present.** There must be a live Polymarket book for the target token.
6. **Spread.** Reject if the Polymarket bid/ask spread exceeds a cap ($\sim 4\text{--}5\text{¢}$), where crossing costs consume the edge.
7. **Book staleness.** Reject if the newest book snapshot is older than 5 s.
8. **Price bounds.** Only trade when the outcome token’s ask is in $[0.25, 0.80]$. As Section 7.3 shows empirically, prices near 0.50 are close to a coin flip.
9. **Dynamic threshold.** A signal-strength-scaled price check.
10. **Position and fill.** The slot must be flat, within risk limits, and the fill attempt itself must succeed.

4.2 Order Execution

Orders are executed as fill-and-kill (FAK) marketable orders, Polymarket’s immediate-or-cancel, which fill at the quoted price or better and cancel any unfilled remainder. A single FAK priced at the second book level automatically sweeps the top level first (it is a better price) and then fills into the second, capturing both levels with one order and no stepping logic. Buys are priced to take displayed liquidity; sells (closes), where exiting matters more than price, are priced more aggressively and post a resting good-till-cancelled remainder if the FAK sweep leaves size unfilled. Polymarket’s taker fee for these markets is modelled per contract as

$$f(p) = p \cdot 0.25 \cdot (p(1-p))^2,$$

which peaks near the middle of the price range and vanishes at the tails; the price-bounds gate is designed partly around it. As Section 7 shows, fees are the largest cost in this strategy.

4.3 Exit Paths

A position is closed by whichever of five conditions fires first:

1. **Stop-loss.** The outcome token's bid has collapsed past a threshold (20–30% of entry).
2. **Reversal.** The Binance underlying itself has reversed through a per-asset threshold, invalidating the thesis at the source.
3. **Convergence.** After a minimum hold (~ 3.5 s), the Polymarket mid has moved through entry in our favour and Binance velocity has exhausted. This is the intended, profitable exit and accounts for 78% of closes.
4. **Timeout.** A per-asset maximum hold (3–27 s) has elapsed.
5. **Emergency.** A 15 s dead-man timer that force-closes regardless of state. As Section 7.3 shows, trades that reach this timer are the strategy's worst cohort.

4.4 Sizing and Risk Limits

Sizing is intentionally small and bounded, appropriate for a strategy whose capacity is unknown and whose fills are uncertain. Orders default to \$100 notional (also the per-order cap), with a hard ceiling of 200 contracts per trade. Exposure is capped at \$500 per asset and \$2,000 in aggregate, with no more than 10 open orders at once. A dual-window rate limiter (a 10-second burst bucket and a 10-minute sustained bucket, mirroring Polymarket's published limits) gates every API call and always reserves tokens for closes, so the system can never rate-limit itself out of exiting a position.

5 Infrastructure

The infrastructure exists to serve one requirement: move a decision from beside Binance to beside Polymarket, and act on it, faster than the Polymarket book reprices. Every design choice below follows from that.

5.1 Two-Node Topology

Because Binance and Polymarket are reached from different continents, detection and execution are split across two nodes (Figure 2). The **Tokyo** node does all Binance-facing work and no execution; the **Dublin** node does all Polymarket-facing work and no detection. They communicate over a single, minimal, connectionless link, so that each node is physically close to the venue it talks to and the expensive intercontinental hop is crossed exactly once, by the smallest message and as early as possible.




Figure 2: Two-node architecture. Tokyo detects Binance moves and transmits a 20-byte UDP signal; Dublin receives, validates against the live Polymarket book, and executes. Each node is co-located with the venue it serves.

5.2 The Wire Protocol

The link between the nodes is a hand-rolled binary protocol carried over UDP. TCP was rejected deliberately: a lost packet on a 120 ms link would trigger a retransmit that arrives long after the edge is gone, so reliability through retransmission is worthless here. Instead, each signal is packed into 20 bytes and sent as **three redundant copies**; the receiver processes the first to arrive and discards duplicates by sequence number. This trades a few extra bytes for resilience to single-packet loss without ever waiting for a round trip. The full frame is given in Table 2.

Table 2: Signal wire format (20-byte payload, big-endian). One frame carries either an OPEN or a CLOSE; bytes 13–14 are repurposed as a reference id on CLOSE.

Offset	Field	Bytes	Description
0	magic	1	0xA7 marks a valid packet
1	seq	2	sequence number (wraps at 65535)
3	copy_id	1	redundant copy index (0, 1, 2)
4	timestamp	8	epoch microseconds
12	flags	1	bit-packed (below)
13	magnitude	1	OPEN: move in bps / CLOSE: ref id high
14	confidence	1	OPEN: score 0–255 / CLOSE: ref id low
15	price	4	Binance price $\times 10,000$
19	checksum	1	XOR of bytes 0–18

The single **flags** byte packs everything needed to route and interpret the signal: a debug bit, a two-bit close reason, a two-bit asset selector, a direction bit, and an action bit. A separate heartbeat frame (magic 0xA8, sent every 2 s) lets Dublin detect if Tokyo goes dark: if heartbeats stop, Dublin knows no signals are coming and existing positions ride out their timeouts safely. When Tokyo sets the debug bit, Dublin forces paper mode and captures extended diagnostics, so a single flag on the sending side switches the whole system between live and instrumented dry-run.

5.3 Node-Internal Design

Both nodes follow the same discipline: the thread doing network I/O never does heavy computation, and the thread doing computation never does I/O, with a single non-blocking queue between them. **Tokyo** runs an asyncio event loop on one core (Binance WebSockets in, status logging) and a worker thread on another (four rolling-window detectors, exit monitoring, signal packing, UDP send); handoff is an $O(1)$ `put_nowait` that can never block the feed. **Dublin** runs all WebSocket feeds (the Polymarket CLOB, the relayed Binance/Chainlink price stream, and the user fill channel) on its asyncio core, and eight worker threads (one per (asset, time-frame) slot) on the other, each owning its own pre-allocated numpy ring buffers so that no worker ever contends with another for its market’s data. Per-slot position state is a small explicit finite state machine (FLAT $\rightarrow$ OPENING $\rightarrow$ LONG/SHORT $\rightarrow$ CLOSING $\rightarrow$ FLAT) locked across the two-phase submit/resolve of every order, which makes concurrent duplicate opens structurally impossible.

5.4 Mock-to-Live Parity

A standalone mock CLOB server replicates Polymarket’s HTTP and WebSocket API: the same order endpoints, the same 200-status-with-error-body convention, the same published rate limits, and, most importantly, order books populated from *Polymarket’s real market data feed*. The mock subscribes to the live Polymarket book over WebSocket and simulates fills by walking the displayed levels. The entire Dublin system talks to the mock or to production through the same code path; going live is a single change of host URL, with no code change. This parity was a deliberate engineering goal, and it succeeded technically. It also, unavoidably, reproduced the market’s blind spot, the subject of Section 8.

5.5 Latency Budget

Figure 3 decomposes the measured end-to-end path. Detection near Binance costs ~ 3 ms; in-node compute in Tokyo is sub-millisecond; the UDP hop to Dublin is ~ 120 ms; Dublin’s gate evaluation and ECDSA order signing add a few milliseconds; and the order reaches Polymarket in ~ 11 ms. The total, ~ 140 – 150 ms, sits at the fast end of the repricing window. The system was fast enough, but the transatlantic segment consumes most of the budget, and no local optimisation can shrink it.




Figure 3: Measured end-to-end latency from a Binance move to a Polymarket order. The ~ 120 ms Tokyo→Dublin transit dominates; all node-internal compute together is a few milliseconds.

6 Deployment

6.1 Instances and Placement

The system ran on two AWS `c6a.large` instances (2 vCPU each, roughly \$52 per month on-demand), one in `ap-northeast-1` (Tokyo) and one in `eu-west-1` (Dublin). “Co-location” here means region and availability-zone placement in the same metro as each venue, with connectivity over AWS’s internal backbone rather than the public internet. That backbone gives lower and, more importantly, lower-jitter latency than a commodity path, without the cost or lead time of true datacenter cross-connects. Two vCPUs matches each node’s two-core design: one core for I/O, one for compute, pinned explicitly. The whole system cost roughly \$100/month to run.

6.2 Measured Latencies

From the deployed nodes, the measured characteristics were approximately 3 ms from Tokyo to Binance, 120 ms from Tokyo to Dublin, and 11 ms from Dublin to Polymarket. These are the numbers behind the budget in Figure 3, and they held stably enough that transit variance was never the binding constraint.

6.3 Operations

Each node logs a compact status line at a fixed interval: Tokyo every 5 s (queue depth, signals sent/dropped, per-asset mid, order-flow imbalance, recent move, book imbalance), Dublin every 15 s (fills, cancels, timeouts, exposure, realised P&L, open orders, rate-limit headroom, heartbeat age). Every state change is written to a dual JSONL-and-parquet trade log, and a post-session analyser computes P&L, win-rate, per-asset breakdown, and exit-reason distribution from the log at shutdown. The debug flag adds per-event timing histograms and full book-state captures at signal time, at zero overhead when disabled.

7 Backtest Performance

The results in this section come from replaying the full strategy (detection, gates, execution logic, fees, and exits) over 13 recorded feed sessions (8,576 trades across roughly 52 hours of order flow). The analysis is confined to the 5-minute markets, which is deliberate rather than a limitation: the large majority of the strategy’s profit came from the 5-minute windows, where the higher window-rotation frequency produces more repricing events per hour than the 15-minute markets. Fills are simulated against the recorded Polymarket book. As Section 8 explains, that fill assumption is what did not hold live; these numbers therefore describe the strategy’s behaviour *given* that displayed liquidity is real, which is the right way to read every figure below.

7.1 Headline Results

The strategy won 61.4% of its 8,576 trades, with a median holding period of 7.6s. The per-trade net P&L distribution (Figure 4) is centred just above zero and positively skewed, consistent with a small, repeatable edge rather than occasional large wins. Exits were dominated by the intended convergence path (78%), with the remainder split across emergency timeouts (16%), opposite-signal closes (4%), and stop-losses (2%).




Figure 4: Per-trade net P&L (cents, clipped to $\pm 25\text{¢}$ for display). Positive mean and median, 61.4% win rate, across 8,576 trades.

7.2 Fees Dominate

The most notable feature of the equity curve (Figure 5) is the gap between gross and net. The strategy earned +51,350¢ gross but paid 28,403¢ in fees (55% of gross), leaving +22,948¢ net. This is the basic arithmetic of a high-turnover strategy on a sub-cent edge: even at Polymarket’s fee schedule, more than half of every unit of gross edge is consumed at the point of execution. It also sets a demanding bar for live execution, since any degradation in fill quality comes directly out of the thin margin that survives fees.




Figure 5: Backtest equity curve over 8,576 chronological trades. Gross P&L (grey) rises steadily; fees (shaded) consume 55% of it, leaving the net curve (blue). The edge is real; the fee drag is large.

7.3 The Edge Has Structure

The edge varies systematically with conditions the strategy can observe at entry (Figure 6), and each relationship has a microstructural reading.

- **Tick rate (panel a).** Win rate falls monotonically from 68% in quiet books (< 30 ticks/s) to a losing 49% in the fastest (> 300 ticks/s). Section 7.4 decomposes the cause.
- **Distance from 0.50 (panel b).** Near 0.50 the outcome is close to a coin flip (51%); the edge strengthens toward the wings (65–68%), where the token price is more sensitive to the underlying and the fair-value decoupling from a Binance move is larger.
- **Hold time (panel c).** P&L is sharply bimodal. The profitable cohorts exit at 2–5 s (+19,774¢) and 8–12 s (+14,243¢); the trades that run to the 12–20 s emergency-timeout band are heavily negative (−10,531¢). This is the same ~ 2 s edge horizon seen in Figure 1: hold past it and the position is systematically underwater.
- **Asset (panel d).** All four assets were net profitable in replay, with SOL and XRP posting the highest win rates ($\sim 65\%$) and BTC/ETH the largest absolute P&L.

Conditional Structure of the Edge — 8,576 backtest trades




Figure 6: Conditional structure of the edge across 8,576 backtest trades. The edge concentrates in quiet books, away from 0.50, at short holding periods, and inverts in each opposite condition.

These are the properties of a healthy microstructure edge: profitable on average, coherent in its drivers, and consistent with the measured repricing dynamics. That is what made the live outcome informative.

7.4 Why the Tick-Rate Effect: Reversal, Not Throughput

The tick-rate relationship (panel a above) is the strongest and most actionable of the four, so its cause is worth isolating. One candidate explanation is throughput: that the system cannot keep up when the Binance feed is fast. The data rules this out. The relationship appears in the *replay* backtest, which processes every recorded tick and does not model the live event queue at all; if it were a throughput artefact, it could not appear here. The cause is therefore in the market, and Figure 7 shows it has two parts.

Why the edge fades in fast markets — reversal + competition, not throughput




Figure 7: Decomposing the tick-rate effect. (a) Win rate (blue) is almost a mirror image of how often the triggering Binance move reverses within one second (red). The moves are not larger at high tick rate (~ 2.2 bps throughout), only more likely to reverse, so reversed and stopped-out exits (amber) climb with it. (b) The Polymarket spread tightens from $\sim 1.9\text{c}$ to $\sim 0.5\text{c}$ as activity rises, evidence of more competitive quoting.

Adverse selection at the source is the dominant part. In quiet books (< 30 ticks/s) the triggering move reverses within a second only 8.5% of the time; in the fastest (> 300 ticks/s) it reverses 28.6% of the time, a $3.4\times$ increase. Since the move magnitude is essentially flat across the range, fast-market signals are more likely to be transient liquidity spikes that snap back, and following momentum into them is adversely selected. That is why reversed and stopped-out exits rise from $\sim 3\%$ to $\sim 14\%$ over the same range.

Competition is the secondary part. As Binance activity rises, the Polymarket book is quoted more tightly (the mean spread falls from $\sim 1.9\text{c}$ to $\sim 0.5\text{c}$) because more market-makers and competing latency arbitrageurs are active, leaving a smaller and shorter-lived mispricing to capture. The practical consequence is that the < 100 ticks/s gate the strategy already applies is well-founded; the data would further support scaling size down, rather than trading flat, above ~ 120 ticks/s.

8 Production Results and Post-Mortem

8.1 What Happened

In paper trading against the mock CLOB, the system performed as Section 7 describes. Signals fired at sensible rates, gates rejected stale and thin conditions, and marketable orders filled at essentially 100%. The mock filled any order the displayed book could satisfy, and the displayed book always appeared deep enough. Live, with real capital, roughly **one order in ten filled**. Same signals, same gates, same latency. Nine times out of ten, the liquidity shown on the book was gone by the time the order arrived.

8.2 Root Cause: Spoofed Liquidity

The displayed depth on the Polymarket book was, to a large extent, not real. Resting orders that made the book look deep were cancelled the instant a real order tried to cross them: classic spoofing, and especially effective in a low-participation venue where a few actors can shape the whole visible book. The strategy's premise was to *take* that displayed liquidity during the repricing window. When the liquidity was a bluff, there was nothing to take: the FAK order arrived, found the quoted size withdrawn, filled a token amount or nothing, and cancelled the rest.

This failure mode was invisible in paper by construction. The mock CLOB simulated fills by walking the *displayed* book levels, the same book the spoofers controlled. It had no model of cancellation-on-arrival because the real book gives no advance signal of it; the cancellation is a *reaction* to the arriving order, and no amount of pre-trade book data contains it. The paper harness was faithful to the displayed market and therefore inherited the same deception. The more realistic the mock was made (consuming displayed liquidity, enforcing real rate limits, using the real feed), the more convincingly it reproduced the illusion. Every favourable number in Section 7 rests on this one false assumption.

8.3 Why Latency Was Not the Problem

It is useful to separate what failed from what did not. The signal was predictive (Figure 1). The backtest was profitable with coherent structure (Figures 5–6). The latency budget was met (Figure 3). The engineering was correct: the system did what it was built to do, on time. The strategy still lost, because a latency arbitrage implicitly assumes that winning the race entitles you to the liquidity at the finish line, and in this venue it did not.

8.4 Adverse Selection, Restated

Seen through a market-microstructure lens, the fills the system *did* get were adversely selected. On the occasions liquidity did not vanish, it often did not vanish because the resting side was willing to trade, which is disproportionately the case when that side is better informed or the move is about to reverse. The strategy was thus filled mainly when being filled was worst, and starved when being filled would have been best. A displayed-book backtest cannot represent this, because it fills every order symmetrically regardless of why the counterparty is or is not there.

9 Lessons and What I Would Do Differently

9.1 Validate Liquidity, Not Just Signal

The main error was validating the signal thoroughly while assuming the execution surface. A corrected process would treat *realisable* liquidity as a first-class unknown to be measured, not inferred from the screen. Concretely: before committing capital, run a campaign of small live probe orders across the target markets and measure the empirical fill rate as a function of displayed size, order aggressiveness, and time-of-day; compare displayed depth against depth that actually transacts in the historical trade prints; and treat any market where realised fills fall far short of displayed depth as untradeable by a liquidity-taking strategy, however attractive the signal.

9.2 Model Cancellation-on-Arrival

A paper harness for a liquidity-taking strategy must include a fill model that can decline. Even a crude one (filling a displayed level only with some probability, or applying a haircut calibrated to observed live fill rates) would have turned the 61%-win-rate backtest into a losing one and prevented the deployment. The point generalises: a backtest that fills against displayed depth is only valid for liquidity-*providing* strategies; for liquidity-takers it is systematically optimistic and must be corrected with a realistic, ideally empirically-calibrated, fill model.

9.3 Respect the Geography

The ~ 120 ms transatlantic hop was the hard floor of this design, and it left almost no margin against a ~ 2 s edge. Where the price source and execution venue are far apart, that transit cost should be the first number computed, and a strategy should proceed only if its edge comfortably survives it, not merely fits inside it.

9.4 Where This Approach Does Work

None of this invalidates cross-venue latency arbitrage as a class. It works when the liquidity being raced toward is firm: deep, professionally market-made venues where displayed size is largely real and cancellation-on-arrival is the exception. The infrastructure built here (the two-node split, the minimal UDP protocol, the lock-free per-asset pipelines, the mock-to-live parity, and the calibration and meta-labeling stack) is sound and reusable. What must change is the target: a venue whose book can be trusted, or a strategy that provides liquidity rather than taking it.

10 Conclusion

This was a complete, deployed, real-money low-latency trading system built around a real and empirically validated inefficiency, and it failed for a reason that none of the engineering it demonstrates could have fixed. The signal predicted short-horizon repricing; the backtest was profitable with coherent conditional structure; the infrastructure delivered orders inside the repricing window; the deployment was cheap, stable, and correct. The strategy still lost money, because the liquidity it was designed to capture was largely illusory, and because the paper harness, faithful to the displayed book, could not detect this until real capital met the real book.

The engineering artefacts retain their value: a template for co-located two-node systems, a compact loss-tolerant signalling protocol, a disciplined lock-free processing model, and a mock-to-live harness with faithful API parity. The strategic lesson is more portable than any of them. In a latency arbitrage, speed is necessary but not sufficient; the binding question is whether the liquidity at the end of the race is real. On Polymarket's short-dated crypto markets, in this period, for a liquidity-taker, it was not. Establishing that with a small, well-instrumented, inexpensive deployment rather than a large one is the useful result here.

A System Reference

A.1 Component Map

Table 3: Principal components by node.

Component	Node	Responsibility
BinanceFeed	Tokyo	Binance depth + trade WebSocket ingest; inter-snapshot book estimation
RollingWindow	Tokyo	60 s per-asset metrics (mid, OFI, VWAP, flow, imbalance, price ring)
LargeMoveDetector	Tokyo	Three-gate signal detection and confidence scoring
ExitMonitor	Tokyo	Tracks active signals; emits CLOSE on reversal/timeout
SignalBridge/Sender	Tokyo	Packs 20-byte frames; sends 3 redundant UDP copies; heartbeat
SignalReceiver	Dublin	UDP ingest; sequence-based dedup; dead-man detection
PolyClobFeed	Dublin	Polymarket CLOB book/trade WebSocket; window rotation
MarketWorker ($\times 8$)	Dublin	Per-(asset, timeframe) event processing over private ring buffers
trade_logic	Dublin	Ten entry gates, five exit paths
OrderEngine/Executor	Dublin	ECDSA signing; FAK/FOK execution; fill reconciliation
PositionManager/FSM	Dublin	Per-slot state machine; two-phase open/close; risk gating
RateLimiter	Dublin	Dual-window API limiter with close reserve

A.2 Backtest Summary

Table 4: Replay backtest summary (paper fill assumption).

Metric	Value
Trades	8,576
Recorded sessions	13 ($\sim$ 52 hours, 5m markets)
Win rate	61.4%
Median hold	7.6 s
Gross P&L	+51,350¢
Fees	28,403¢ (55% of gross)
Net P&L	+22,948¢
Convergence exits	78%
Emergency / opp. / stop	16% / 4% / 2%
Repricing half-life	$\sim$ 786 ms (median)
Live fill rate	$\sim$10% (vs $\sim$100% paper)

A.3 Key Parameters

Table 5: Selected system parameters.

Parameter	Value	Where
Assets	BTC, ETH, SOL, XRP	both
Timeframes	5 min, 15 min	Dublin
Detector lookback	25–150 ms	Tokyo
Signal packet size	20 bytes ($\times$ 3 copies)	wire
Heartbeat interval	2 s	wire
Signal-to-wire budget	50–200 μ s	Tokyo
Max signal age gate	500 ms	Dublin
Price-bounds gate	[0.25, 0.80]	Dublin
Default order size	\$100 (cap \$100)	Dublin
Exposure caps	\$500 / asset, \$2,000 total	Dublin
Emergency dead-man	15 s	Dublin
Instances	2 $\times$ c6a.large (2 vCPU)	AWS
Regions	ap-northeast-1, eu-west-1	AWS
End-to-end latency	$\sim$ 140–150 ms	measured
Monthly cost	$\sim$ \$100	AWS